

Report – CSCI-B 651

Mini GPT-Style Language Model

Meghana Nanuvala
May 2, 2025

Environment Setup

The development environment was set up using Google Colab with GPU acceleration enabled. The availability of the GPU was verified using `torch.cuda.is_available()`.

GPT-style language models require substantial computational resources. Using Colab's free GPU significantly accelerates training. This makes it possible to train deep learning models within assignment constraints. Due to limitations in Google Colab for extended training sessions, Kaggle's dual-GPU environment was used for the full-scale model training.

1. Data Collection and Preprocessing

Dataset Source and Justification for Size

The dataset used was the English Wikipedia dump available via the Hugging Face *datasets* library. It was loaded and a subset of 100,000 full-length articles was selected,

```
1 from datasets import load_dataset
2 import re
3
4 raw_dataset = load_dataset("wikimedia/wikipedia", "20231101.en", split="train")
5 subset = raw_dataset.select(range(100000))
6
```

Justification for Size:

The recommended article size in the assignment was ~100,000 articles. This size was chosen to balance training quality and computational feasibility, especially when using Kaggle's 2x Tesla T4 GPU setup. Empirical testing showed that training on 100,000 articles would generate approximately 861,981 sequences of 128 tokens each, which is large enough for the model to learn complex structures while remaining trainable in under 5 hours.

Preprocessing Pipeline

The raw text content in each article often contains Wikipedia-specific markup, metadata, and formatting. These are inappropriate for model training. Therefore, a multi-step text cleaning process was implemented using regular expressions and rule-based filters.

The following steps were applied to each article's text field,

1. Removal of Wiki Markup

All double-bracketed internal links were removed: `[[]]`. Section headings such as `== Heading ==` were stripped. Templates like `{{Infobox}}` and table-like structures were discarded.

2. Elimination of Non-Textual Content

HTML/XML tags (e.g., <ref>,
) were removed. URLs and file/media references such as https://..., File:..., and Image:... were excluded.

3. Whitespace and Sentence Cleanup

Multiple line breaks and tab characters were replaced with single spaces. Repeated punctuation or formatting symbols were normalized. Extra spaces were trimmed to standardize spacing.

4. Splitting into Paragraphs

Articles were split into individual paragraphs using newline or period-based segmentation. Each paragraph was treated as one sample and saved as a separate line.

5. Filtering for Minimum Length

Paragraphs with fewer than 5 words or lacking alphabetic characters were excluded to maintain linguistic quality.

Format and Storage

After preprocessing, each paragraph was written as a single line in a plain-text file:

```
22 with open("cleaned_wiki_100k.txt", "w", encoding="utf-8") as f:
23     for paragraph in cleaned_texts:
24         if paragraph:
25             f.write(paragraph.replace("\n", " ").strip() + "\n")
26
```

This resulted in a file compatible with both tokenizer training and later model input, following the structure:

- One paragraph per line
- Plain .txt format
- Total size: ~1.2 million cleaned lines from the original 100,000 articles

Example Entries from Processed Dataset

Here are representative examples of cleaned paragraphs after preprocessing:

```
Allan Dwan (born Joseph Aloysius Dwan; April 3, 1885 – December 28, 1981) was a pioneering Canadian-born American motion pict
Algeria, officially the People's Democratic Republic of Algeria, is a country in North Africa. Algeria is bordered to the nor
This is a list of characters in Ayn Rand's 1957 novel Atlas Shrugged. Major characters The following are major characters fr
Anthropology is the scientific study of humanity, concerned with human behavior, human biology, cultures, societies, and ling
Agricultural science (or agriscience for short) is a broad multidisciplinary field of biology that encompasses the parts of e
Anarchism is a political philosophy and movement that is skeptical of all justifications for authority and seeks to ab
Albedo (; ) is the fraction of sunlight that is diffusely reflected by a body. It is measured on a scale from 0 (corre
A, or a, is the first letter and the first vowel of the Latin alphabet, used in the modern English alphabet, the alpha
Alabama () is a state in the Southeastern region of the United States, bordered by Tennessee to the north; Georgia to
In Greek mythology, Achilles ( ) or Achilleus ( ) was a hero of the Trojan War who was known as being the greatest of a
```

These samples show well-formed English sentences free of markup. This ensures the model can focus on learning linguistic structure rather than parsing noisy symbols.

2. Tokenizer Training

Methodology

Used the tokenizers library by Hugging Face to train a BPE tokenizer from scratch.

Training Setup:

- **Tokenizer Type** : BPE (tokenizers.models.BPE)
- **Pre-tokenizer** : Whitespace (splits on spaces)
- **Trainer Type** : BpeTrainer with vocab size 20,000
- **Dataset** : cleaned_wiki_100k.txt (1 line = 1 paragraph)
- **Special Tokens** :
 - <pad> — padding for fixed-length sequences
 - <unk> — unknown token fallback
 - <s> — beginning-of-sequence marker (used during generation)
 - </s> — optional end-of-sequence marker (not used during generation)

```
3 from tokenizers import Tokenizer, models, trainers, pre_tokenizers, decoders, processors
4 |
5 tokenizer = Tokenizer(models.BPE()) # Initialize tokenizer
6
7 tokenizer.pre_tokenizer = pre_tokenizers.Whitespace() # Pre-tokenizer (splits on whitespace and pu
8
9 special_tokens = ["<pad>", "<unk>", "<s>", "</s>"] # Special tokens
10
11 # Trainer
12 trainer = trainers.BpeTrainer(
13     vocab_size=20000,
14     special_tokens=special_tokens
15 )
16
17 tokenizer.train(files=["cleaned_wiki_100k.txt"], trainer=trainer) # Train from file
18
19 tokenizer.save("bpe_tokenizer_100k.json")
```

Parameter	Choice Made	Justification
Tokenizer Type	BPE	Efficient balance between word-level and character-level representations.
Vocab Size	20,000	Large enough to encode subwords with precision, while keeping the model compact.
Special Tokens	<pad>, <unk>, <s>	Required for padding, unknown fallback, and text generation start markers.
End Token </s>	Included but unused	Kept for possible compatibility, not required in autoregressive generation.

Sample Tokenization Results

Given the cleaned input sentence the output token IDs and the corresponding tokens are,

```
Input Text: Anarchism is a political philosophy.  
Token IDs: [6890, 7499, 7286, 6693, 68, 8374, 13165, 17]  
Tokens: ['An', 'arch', 'ism', 'is', 'a', 'political', 'philosophy', '.']  
Vocabulary Size: 20000
```

The resulting vocabulary contained exactly 20,000 tokens which including subword units such as:

['re', 'tion', 'ization', 'ent', 'pre', 'theo', 'Ein', 'stein', 'ology', 'AI', 'para', ...]

This subword-level vocabulary enables better generalization to rare and unseen words by decomposing them into known subunits. It also has a compact model size and embedding table.

Why Use BPE with GPT Models

GPT-style models use an autoregressive language modeling objective that requires the ability to:

- Encode sequences of arbitrary length efficiently.
- Represent rare, compound, or unknown words.
- Minimize the number of tokens per input.

Byte-Pair Encoding (BPE) is widely used in GPT-style models because it:

- Produces stable, deterministic subword splits.
- Reduces vocabulary explosion.
- Maintains semantic interpretability better than byte-level tokenization.

Using a custom BPE tokenizer aligned to the training data ensures:

- Better token coverage.
- Fewer <unk> tokens during inference.
- Lower perplexity and better generalization.

3. Dataset Preparation

Methodology

Step 1: Load Tokenizer and Cleaned Text

The trained BPE tokenizer (bpe_tokenizer_100k.json) was loaded using the tokenizers library, and the cleaned paragraph-level dataset (cleaned_wiki_100k.txt) was read line by line.

```
9 # Load tokenizer  
10 from tokenizers import Tokenizer  
11 tokenizer = Tokenizer.from_file("bpe_tokenizer_100k.json")  
12  
13 # Load raw text lines  
14 with open("cleaned_wiki_100k.txt", "r", encoding="utf-8") as f:  
15     lines = [line.strip() for line in f if line.strip()]  
16
```

Step 2: Tokenization and BOS Injection

Each line was tokenized and prepended with the <s> token. Token IDs were collected sequentially.

```
5 SEQ_LEN = 128
6 PAD_TOKEN_ID = tokenizer.token_to_id("<pad>")
7 BOS_TOKEN_ID = tokenizer.token_to_id("<s>")
8
18 encoded_lines = []
19 for line in lines:
20     tokens = tokenizer.encode("<s> " + line).ids # add BOS token
21     encoded_lines.extend(tokens + [PAD_TOKEN_ID]) # pad between sequences
22
```

Step 3: Chunking into Fixed-Length Sequences

The continuous token stream was split into non-overlapping chunks of 128 tokens. Only full-length sequences were kept to ensure shape consistency.

```
23 # Split into chunks of SEQ_LEN
24 def chunk_sequence(data, seq_len):
25     return [data[i:i+seq_len] for i in range(0, len(data), seq_len) if len(data[i:i+seq_len]) == seq_len]
26
27 chunks = chunk_sequence(encoded_lines, SEQ_LEN)
28
```

Step 4: Tensor Creation for Training

The token chunks were converted into PyTorch tensors for use in Keras (later converted to NumPy). Labels were created by shifting the input tensor one position to the left to match the CLM objective (predicting the next token).

```
30 inputs = torch.tensor(chunks)
31 labels = inputs.clone()
32
33 labels = torch.roll(labels, shifts=-1, dims=1) # Shift labels to the right by one
34
35 print("Input Example (token IDs):", inputs[0])
36 print("Label Example (token IDs):", labels[0])
37
38 torch.save(inputs, "inputs_100k.pt")
39 torch.save(labels, "labels_100k.pt")
40
41 print("Saved inputs.pt and labels.pt successfully.")
42
```

Why This Preparation Is Necessary

GPT-style models expect training data in the form of:

- **Fixed-length sequences** : to enable batch training and GPU acceleration.
- **Right-padded where needed** : to ensure uniformity across batches.
- **Autoregressive format** : where each position predicts the next token.

By creating inputs and labels with this structure the model can efficiently learn causal dependencies. The model uses masked cross-entropy loss that ignores padding tokens for efficient learning.

Example from Prepared Dataset

```
Input Example (token IDs): tensor([ 2, 6890, 7499, 7286, 6693, 68, 8374, 13165, 6696, 9545,
6770, 6693, 86, 11274, 6849, 6695, 6762, 7875, 11771, 6718,
11859, 6696, 7916, 6843, 6701, 14520, 6780, 6685, 11241, 6698,
11579, 9509, 6726, 11645, 6976, 6682, 70, 6702, 6696, 18368,
17866, 15, 10666, 7329, 10030, 16, 8920, 15, 6696, 9351,
7286, 17, 6890, 7499, 7286, 12749, 6950, 6718, 6685, 14113,
6695, 6685, 7396, 6750, 8099, 13167, 17809, 6696, 18958, 6821,
8640, 18380, 17, 7020, 68, 16466, 7689, 16, 9897, 9545,
15, 6988, 9455, 6695, 19669, 7286, 6693, 8068, 6684, 6685,
8804, 6685, 6705, 7689, 6695, 6685, 8374, 18681, 15, 8805,
8518, 6694, 6685, 6987, 7595, 9681, 9897, 6695, 6685, 8755,
6754, 9545, 11, 6987, 7595, 9681, 8755, 7286, 6956, 9409,
6918, 6970, 9687, 6681, 17809, 8399, 11743, 18368])
Label Example (token IDs): tensor([ 6890, 7499, 7286, 6693, 68, 8374, 13165, 6696, 9545, 6770,
6693, 86, 11274, 6849, 6695, 6762, 7875, 11771, 6718, 11859,
6696, 7916, 6843, 6701, 14520, 6780, 6685, 11241, 6698, 11579,
9509, 6726, 11645, 6976, 6682, 70, 6702, 6696, 18368, 17866,
15, 10666, 7329, 10030, 16, 8920, 15, 6696, 9351, 7286,
17, 6890, 7499, 7286, 12749, 6950, 6718, 6685, 14113, 6695,
6685, 7396, 6750, 8099, 13167, 17809, 6696, 18958, 6821, 8640,
18380, 17, 7020, 68, 16466, 7689, 16, 9897, 9545, 15,
6988, 9455, 6695, 19669, 7286, 6693, 8068, 6684, 6685, 8804,
6685, 6705, 7689, 6695, 6685, 8374, 18681, 15, 8805, 8518,
6694, 6685, 6987, 7595, 9681, 9897, 6695, 6685, 8755, 6754,
9545, 11, 6987, 7595, 9681, 8755, 7286, 6956, 9409, 6918,
6970, 9687, 6681, 17809, 8399, 11743, 18368, 2])
Saved inputs.pt and labels.pt successfully.
```

- The input begins with the BOS token (2).
- The label is shifted by one to match the next token prediction target.
- Both tensors have the shape: (861981, 128)

4. Model Architecture and Training

This step involved designing and training a GPT-style causal language model, defined as a unidirectional Transformer architecture composed entirely of decoder blocks. The model is trained using causal language modeling (CLM). Each token is predicted given all tokens before it, with padding properly masked during loss computation.

Implementation Overview

The model was implemented using Keras and the keras-nlp library. This library provides efficient and modular transformer components suitable for GPT-like architectures. To accelerate training multi-GPU training was enabled using `tf.distribute.MirroredStrategy`.

Model Configuration and Hyperparameters

Hyperparameter	Value
Vocabulary Size	20,000 (from trained BPE)
Sequence Length	128 tokens
Embedding Dimension	256
Transformer Layers	4
Attention Heads	4
Feedforward Dim	1024

Dropout	Not applied (default 0)
Total Parameters	~13.45M
Optimizer	Adam (lr=1e-4, clipnorm=1.0)
Loss Function	SparseCategoricalCrossentropy (from logits)
Metric	Perplexity (mask_token_id=0)
Epochs	2
Batch Size	32 × 2 GPUs = 64

Model Architecture

```

31 strategy = tf.distribute.MirroredStrategy()
32 print(f"Number of GPUs: {strategy.num_replicas_in_sync}")
33 GLOBAL_BATCH_SIZE = BATCH_SIZE * strategy.num_replicas_in_sync
34
35 # Dataset Preparation
36 train_dataset = tf.data.Dataset.from_tensor_slices((inputs, labels))
37 train_dataset = train_dataset.shuffle(1000).batch(GLOBAL_BATCH_SIZE).prefetch(tf.data.AUTOTUNE)
38
39 # Model Definition and Training
40 with strategy.scope():
41     inputs_ = keras.layers.Input(shape=(SEQ_LEN,), dtype="int32")
42     x = keras.layers.Masking(mask_value=PAD_TOKEN_ID)(inputs_)
43
44     x = keras_nlp.layers.TokenAndPositionEmbedding(
45         vocabulary_size=VOCAB_SIZE,
46         sequence_length=SEQ_LEN,
47         embedding_dim=EMBED_DIM,
48         mask_zero=True
49     )(x)
50
51     for _ in range(NUM_LAYERS):
52         x = keras_nlp.layers.TransformerDecoder(
53             num_heads=NUM_HEADS,
54             intermediate_dim=FEED_FORWARD_DIM,
55         )(x)
56
57     outputs = keras.layers.Dense(VOCAB_SIZE)(x)
58     model = keras.Model(inputs=inputs_, outputs=outputs)
59
60     loss_fn = keras.losses.SparseCategoricalCrossentropy(from_logits=True)
61     perplexity = keras_nlp.metrics.Perplexity(from_logits=True, mask_token_id=PAD_TOKEN_ID)
62     optimizer = keras.optimizers.Adam(learning_rate=0.0001, clipnorm=1.0)
63
64     model.compile(optimizer=optimizer, loss=loss_fn, metrics=[perplexity])
65
66     model.summary()

```

Component	Function
Input(shape=(128,))	Accepts token ID sequences of length 128.
Masking(mask_value=0)	Ensures that padded tokens (<pad> = 0) are ignored during training.
TokenAndPositionEmbedding	Adds learned embeddings for token ID and position.
TransformerDecoder (×4)	Unidirectional self-attention blocks with feedforward layers.
Dense(vocab_size)	Final classification layer over vocabulary to predict next token.
SparseCategoricalCrossentropy	Suitable for integer targets; combined with logits output.

Component	Function
Perplexity Metric	Measures the model's confidence in predicting the next token.
Adam Optimizer	Adaptive learning with gradient clipping for stability.

Training Logs (Kaggle, 2x T4 GPUs)

```
Epoch 1/2
13469/13469 ————— 2153s 158ms/step - loss: 5.9798 - perplexity: 562.1618
Epoch 2/2
13469/13469 ————— 2145s 159ms/step - loss: 4.6121 - perplexity: 100.7642
```

Epoch	Loss	Perplexity
1	5.9798	562.1618
2	4.6121	100.7642

Perplexity dropped by ~82%, demonstrating strong learning progress over 2 epochs. The model successfully learned token dependencies and sentence structure from 861,981 input sequences.

Model Summary (Partial Output)

inputs and labels loaded and converted to NumPy.
Shape: (861981, 128) (861981, 128)
Number of GPUs: 1
Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 128)	0
masking (Masking)	(None, 128)	0
token_and_position_embedding (TokenAndPositionEmbedding)	(None, 128, 256)	5,152,768
transformer_decoder (TransformerDecoder)	(None, 128, 256)	789,760
transformer_decoder_1 (TransformerDecoder)	(None, 128, 256)	789,760
transformer_decoder_2 (TransformerDecoder)	(None, 128, 256)	789,760
transformer_decoder_3 (TransformerDecoder)	(None, 128, 256)	789,760
dense (Dense)	(None, 128, 20000)	5,140,000

Total params: 13,451,808 (51.31 MB)
Trainable params: 13,451,808 (51.31 MB)

Justification of Design Decisions

The used 4 decoder layers are sufficient for modeling dependencies without overloading memory. 256 embedding dimensions provides a balanced expressiveness and a training efficiency. The token sequence length of 128 is chosen since it is long enough for paragraph-level modeling. To enable faster training and higher batch sizes Multi-GPU via MirroredStrategy were used.

5. Evaluation and Text Generation

Quantitative Evaluation:

Final Validation Metrics (Epoch 2):

- **Loss:** 4.6121
- **Perplexity:** 100.76

Interpretation:

The perplexity score represents the model's average uncertainty when predicting the next token. A lower value indicates more confident predictions.

By considering:

- **Vocabulary size:** 20,000
- **Model size:** 13.45 million parameters
- **Training time:** 2 epochs
- **Training data:** 861,981 sequences from 100,000 articles

A perplexity score of ~100 is reasonable and expected. It shows the model has successfully learned general sentence structure, common word patterns, and token transitions from the dataset.

Text Generation Setup

Text generation was performed using greedy decoding. A trained tokenizer (bpe_tokenizer_100k.json) was used to encode prompts. The model was autoregressively generating tokens until a maximum sequence length was reached.

- **Sequence length:** 128 tokens
- **Generation method:** Greedy decoding
- **Decoding start:** <s> token + optional prompt
- **Tokenizer:** Trained Byte-Pair Encoding tokenizer with 20K vocab

Generated Examples

Prompt: "The future of AI "

Output: *The future of AI - I is a list of people who are the most important of the world ' s history. The first of the world is the most important of the world ' s history. The world ' s history is the most important of the world ' s history.*

Prompt: "Climate change is "

Output: *Climate change is a climate change in the climate change in the climate change. The climate change is a climate change in the climate change in the climate change. The climate change is a climate change in the climate change.*

Prompt: "Albert Einstein was"

Output: *Albert Einstein was a German-language film maker. He was born in the Dutch province of North Holland. He was a member of the German Academy of Arts and Sciences. He was born in the Dutch province of North Holland.*

```
◆ Example 1 – Prompt: The future of AI
The future of AI - I is a list of people who are the most important of the world ' s history . The first of the world is the most important of th

◆ Example 2 – Prompt: Climate change is
Climate change is a climate change in the climate change in the climate change . The climate change is a climate change in the climate change . Th

◆ Example 3 – Prompt: Albert Einstein was
Albert Ein stein was a German - language film maker . He was born in the Dutch province of North Holland . He was a member of the German Academy o
```

Qualitative Evaluation:

Evaluation Aspect	Assessment
Syntactic Fluency	Sentences are grammatically correct and punctuated well.
Lexical Diversity	Improved compared to earlier models; model accesses rich vocabulary.
Structural Consistency	Sentence forms are consistent with natural English syntax.
Semantic Coherence	Weak — ideas loop or repeat; content lacks logical development.
Repetition	Present in outputs, especially with factual phrases ("climate change...").
Named Entity Use	Sometimes plausible (e.g., Einstein + German Academy) but likely memorized.

6. Critical Analysis and Justification

Vocabulary Size and Tokenization Trade-Offs

A Byte-Pair Encoding (BPE) tokenizer with a vocabulary size of 20,000 was selected to balance expressiveness and computational efficiency. Larger vocabularies such as the 50K+ used in large-scale GPT models therefore can better capture full words and rare entities. In contrast, the 20K size enables subword coverage for uncommon or compound words. This reduces the frequency of <unk> tokens, and remains practical for small-scale training environments like Kaggle.

The choice of BPE over character or word level tokenization was made because it offers:

- Better compression than word-level tokenization.
- More semantic clarity than character-level tokenization.
- Faster training and inference due to shorter sequence lengths.

Model Size and Resource Constraints

The model was limited to 4 decoder layers (~13.45M parameters) due to GPU constraints. While sufficient for capturing basic grammatical structures, this configuration restricts the model's ability to learn long-range semantic dependencies. As observed in generated outputs while sentences were fluent and grammatically correct but deeper contextual understanding was lacking.

A deeper model (6-8 layers) with larger embedding dimensions and attention heads would enhance its ability to capture nuanced relationships between tokens. Also will maintain coherence over longer sequences. However such enhancements require extended training time and memory beyond what was feasible in this assignment's scope.

Dataset Quality and Impact

The dataset comprising 100,000 cleaned English Wikipedia articles provided a high-quality formal language rich in real-world entities.

This allowed the model to:

- Learn sentence boundaries and punctuation.
- Recognize named entities (e.g., “Albert Einstein,” “North Holland”).
- Produce fluent and well-formed output.

However, Wikipedia’s templated structure, often biographical or encyclopedic, encouraged repetitive and factoid-style generations. This limited the diversity of generated topics and reduced creativity in open-ended prompts.

Coherence and Output Limitations

Despite syntactic fluency outputs often suffered from:

- Repetition loops.
- Weak narrative progression.
- Overuse of memorized phrase structures.

These shortcomings are expected in GPT models trained on shallow datasets for a few epochs. Improved coherence requires either:

- More training epochs.
- Richer or more diverse training data.
- Architectural enhancements like relative positional encoding or longer context windows.

Recommendations for Future Work

To enhance both semantic coherence and generation diversity the following improvements can be considered:

- Train for longer epochs to better internalize token relationships.
- Use advanced decoding strategies such as Top-k or nucleus (top-p) sampling to avoid repetition.
- Add dropout and normalization layers which will improve generalization and prevent overfitting to frequent token paths.
- Train on broader corpora which includes narrative, dialogue, or QA datasets to diversify patterns.
- Increasing model depth would support richer abstraction.

Overall Justification

Every component of the pipeline was deliberately selected to operate within the computational constraints while still demonstrating the fundamental capabilities of autoregressive text generation. The model achieved a validation perplexity of 100.76, supporting the claim that it learned meaningful patterns from the training data.

Conclusion

This project validated the feasibility of building a GPT-style language model from scratch using a real-world corpus, efficient tokenization, and modern Transformer components. While future improvements are necessary for achieving higher semantic coherence, the current model performs well for its size and scope. Current model was able to achieve successful end-to-end language modeling on a moderately sized dataset.

References

Hugging Face Datasets wikimedia/wikipedia Dataset. Available at:

<https://huggingface.co/datasets/wikimedia/wikipedia>

Hugging Face Tokenizers Library Available at: <https://github.com/huggingface/tokenizers>

KerasNLP Natural Language Processing with Keras. Available at:

<https://github.com/keras-team/keras-nlp>

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N & Polosukhin, I. (2017). Attention is All You Need. In Advances in Neural Information Processing Systems (NeurIPS). Available at: <https://arxiv.org/abs/1706.03762>

OpenAI (2020). Language Models are Few-Shot Learners. Available at:

<https://arxiv.org/abs/2005.14165>

TensorFlow Documentation MirroredStrategy. Available at:

https://www.tensorflow.org/api_docs/python/tf/distribute/MirroredStrategy